Radial basis function networks

Exercise T7.1:   Multi-class classification: simple methods   (tutorial)

(a) Describe how a K nearest neighbor classifier predicts the class of previously unseen inputs?

(b) A "Parzen window" classifier extends the *electoral committee* approach of kNN. How are the different votes *weighted*?

Exercise T7.2:   Radial basis function networks   (tutorial)

(a) Describe and discuss the *general architecture* of an RBF-network.

(b) Describe and discuss the *two-step learning procedure* for RBF-networks with K basis functions.

(c) In which cases of *regression* or *classification* outperform *MLPs* RBF networks significantly? What are the advantages of RBF networks? In which situations are they preferable to MLP?

Exercise H7.1:   Training data   (homework, 1 point)

Create a sample of $p = 120$ training patterns $\{\underline{\mathbf{x}}^{(\alpha)}, y_T^{(\alpha)}\}$, $\alpha = 1, \ldots, p$. The input values $\underline{\mathbf{x}}^{(\alpha)} \in \mathbb{R}^2$ should be drawn from a mixture of Gaussians with centers in an XOR-configuration according to the following scheme:

- Generate 60 samples from each of the following two conditional distributions:

$$p(\underline{\mathbf{x}}|y = 0) := \tfrac{1}{2}[\mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_1, \sigma^2) + \mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_2, \sigma^2)],$$
$$p(\underline{\mathbf{x}}|y = 1) := \tfrac{1}{2}[\mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_3, \sigma^2) + \mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_4, \sigma^2)],$$

  with $\underline{\boldsymbol{\mu}}_1 = (0,1)^\top, \underline{\boldsymbol{\mu}}_2 = (1,0)^\top, \underline{\boldsymbol{\mu}}_3 = (0,0)^\top, \underline{\boldsymbol{\mu}}_4 = (1,1)^\top$ and a variance of $\sigma^2 = 0.1$. To sample from one of the two mixture variables you can (i) draw with probability 0.5 whether you draw from the density in the left summand or from the density in the right summand, and (ii) sample then from that normal distribution yielding one $\underline{\mathbf{x}}^{(\alpha)}$. Note that $\mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}, \sigma^2)$ is the probability density of a multivariate normal distribution of a vector $\underline{\mathbf{x}}$, where $\underline{\boldsymbol{\mu}}$ is the mean vector and σ^2 the variance (same for all components).

- The corresponding target values $y_T^{(\alpha)} \in \{0, 1\}$ describe the assignment to the two classes and indicate from which distribution [$p(\underline{\mathbf{x}}|y = 0)$ vs. $p(\underline{\mathbf{x}}|y = 1)$] the data point was drawn.

(a) (1 point) Plot the resulting 120 input samples $\underline{\mathbf{x}}^{(\alpha)}$ in a scatter plot, in which the markers and/or colors represent the corresponding samples' labels $y_T^{(\alpha)}$.

Exercise H7.2: k nearest neighbors (kNN)   (homework, 2 points)

Build a kNN classifier that classifies new data (*query points*) by voting of the K nearest neighbors from the training set. Thus the *electoral committee* is selected from the training patterns according their Euclidean distance to the query point. The predicted class is determined by the target values of the majority of those K nearest patterns.

(a) (2 points) Plot the training patterns and the decision boundary (e.g. using a contour plot or a high-resolution image of equidistant query points) in input space for $K = 1, 3, 5$. What do you observe?

Exercise H7.3:   "Parzen window" classifier   (homework, 3 points)

This classifier implements a *weighted voting scheme*. All training points (not only the K nearest ones) make a vote for the query point but their vote is weighted by a *Parzen window* (or *kernel function*) depending on the distance between the training samples $\mathbf{x}^{(\alpha)}$ and the query point $\underline{\mathbf{x}}$. The Gaussian window function based on Euclidean norm $\| \cdot \|$ is:

$$\kappa(\underline{\mathbf{x}}, \mathbf{x}^{(\alpha)}) = \exp\left(-\frac{1}{2\sigma_\kappa^2} \|\mathbf{x} - \mathbf{x}^{(\alpha)}\|^2 \right).$$

(a) (2 points) Plot the training patterns and the decision boundary (e.g. using a contour plot or a high-resolution image of equidistant query points) in input space for Gaussian window functions parameterized with the variances $\sigma_\kappa^2 = 0.5, 0.1$ and 0.01.

(b) (1 point) Rerun kNN and Parzen-window classification after adding 60 more data points from a third class centered on $\underline{\boldsymbol{\mu}}_3 = (0.5, 0.5)^\top$ with variance $\tilde{\sigma}^2 = 0.05$. Plot the classification boundaries as above and compare them with your previous results.

Exercise H7.4:   RBF networks   (homework, 4 points)

Similar to the Parzen window, RBF networks classify data according to a weighted vote, but the voting committee now consists of $K \ll p$ "representatives" instead of all p data points. These representatives do not have to be previously seen data points and can be "prototypes" $\underline{\mathbf{t}}_i \in \mathbb{R}^2$ derived from the training data via K-means clustering. Construct a RBF-net as follows:

- Determine the K representatives $\underline{\mathbf{t}}_i$ via K-*means* clustering (you can implement the online-algorithm described in the lecture notes or use available packages).

- For a given weight vector $\underline{\mathbf{w}} \in \mathbb{R}^{K+1}$, the predicted classification for a query point $\underline{\mathbf{x}}$ is:

$$y(\underline{\mathbf{x}}) = \widetilde{\text{sign}}(\mathbf{w}^\top \underline{\phi}(\underline{\mathbf{x}})),$$

where $\underline{\phi}(\underline{\mathbf{x}}) := \begin{pmatrix} 1 \\ \phi_1(\underline{\mathbf{x}}) \\ \dots \\ \phi_K(\underline{\mathbf{x}}) \end{pmatrix}$ is a $(k+1)$-dimensional vector containing the bias and the basis function values $\phi_i(\underline{\mathbf{x}}) = \kappa(\underline{\mathbf{x}}, \underline{\mathbf{t}}_i)$ with κ from the previous exercise (Gaussian radial basis functions). Here again we use the slightly modified sign function

$$\widetilde{\text{sign}}(h) = \begin{cases} 1 \text{ for } h \geq 0, \\ 0 \text{ otherwise.} \end{cases}$$

(instead of the usual +1/-1 sign function).

- Determine the weight vector as: $\underline{\mathbf{w}} = \left(\underline{\boldsymbol{\Phi}}\,\underline{\boldsymbol{\Phi}}^{\top}\right)^{-1}\underline{\boldsymbol{\Phi}}\,\underline{\mathbf{y}}^{T}$ where $\underline{\mathbf{y}} := (y_T^{(1)}, \ldots, y_T^{(p)}) \in \mathbb{R}^{1,p}$ is the vector of target values and

$$\underline{\boldsymbol{\Phi}} := \left(\underline{\phi}(\underline{\mathbf{x}}^{(1)}), \ldots, \underline{\phi}(\underline{\mathbf{x}}^{(p)})\right) \in \mathbb{R}^{K+1,p}$$

(a) (2 points) Plot the decision boundaries together with the training patterns and locations of the representatives for $K \in \{2, 3, 4\}$. Do this for two different (reasonable) kernel widths σ_κ of the radial basis functions ϕ_i, yielding a total of six plots.

(b) (2 points) Construct a RBF-network with 2 RBFs and fix the centers to $\underline{\mathbf{t}}_1 = (0, 0)^T$ and $\underline{\mathbf{t}}_2 = (1, 1)^T$ (i.e., skip K-means clustering). For $\sigma_\kappa = 0.45$, make a scatter plot of the data in the space of RBF-activations, i.e. for each data point α plot $\phi_1(\underline{\mathbf{x}}^{(\alpha)})$ vs. $\phi_2(\underline{\mathbf{x}}^{(\alpha)})$ and indicate their class-assignment via their color. Plot also the predicted labels after training in a similar second plot. Feel free to reduce the data-variance σ (e.g. to 0.2) to make the cluster-structure more prominent.

Total 10 points.